



Bridge of Life
Education

SOC Design

Kernel IO Interface

Topics

- Overview
- Block Level Protocol
 - ap_ctrl_hs
 - ap_ctrl_chain
 - ap_ctrl_none (datflow) - advanced
- Port Level Protocol
 - s_axilite
 - m_axi - axi master
 - axis - axi stream
- Pragma & Examples

Overview

Mapping of Key Attributes of C Code

Function: design hierarchy, mapped to MODULE

Arguments : mapped to Input/output interface of the hardware

Types: All variables are of a defined type, influence the area and the performance

Loops: impact on area and performance, HLS opt with Directive Pragma

Control flow: Control logic

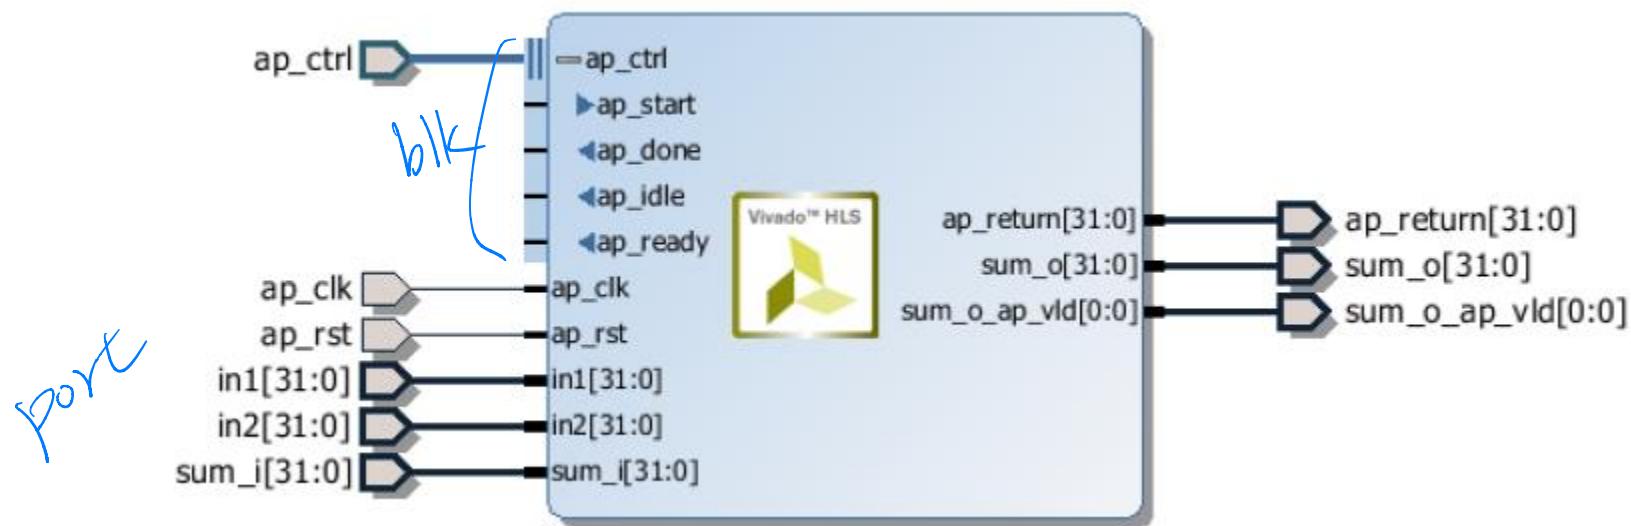
Arrays: impact the device area, and performance bottleneck

Operators: Function unit. Allocation/Scheduling (Sharing) to meet performance and area

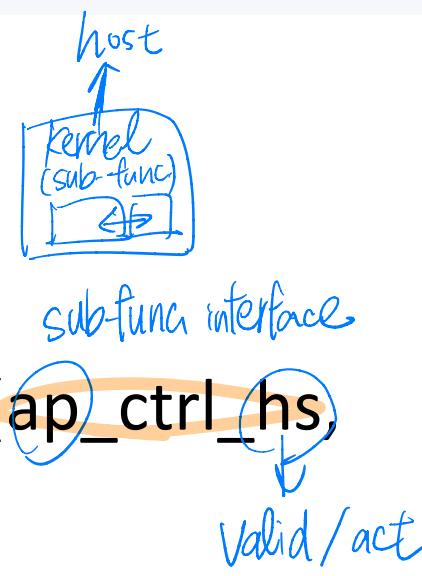
```
46 #include "fir.h"
47
48 void fir (
49     data_t *y,
50     coef_t c [IN],
51     data_t x
52 ) {
53
54
55
56
57     int i;
58
59     acc=0;
60     Shift_Accum_Loop: for (i=N-1;i>=0;i--) {
61         if (i==0) {
62             shift_reg[0]=x;
63             data = x;
64         } else {
65             shift_reg[i]=shift_reg[i-1];
66             data = shift_reg[i];
67         }
68         acc+=data*c[i];
69         *y=acc;
70     }
71 }
```

Overview

```
#include "sum_io.h"
dout_t sum_io (    din_t      in1,
                    din_t      in2,
                    dio_t      *sum) {
    dout_t temp;
    *sum = in1 + in2 + *sum;
    temp = in1 + in2;
    return temp;
}
```

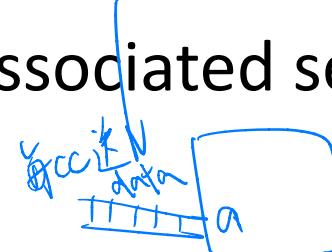


Top Function Communicates with Host

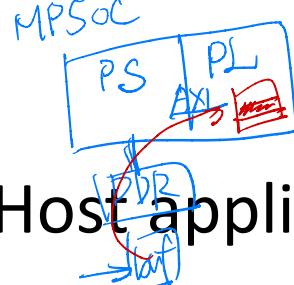


only block func.

- **Block Protocol**: specified on the s_axilite interface (ap_ctrl_hs, ap_ctrl_chain, ap_ctrl_none)
- **Port Protocol**: AXI4-Master, AXI-Lite , AXI4-Stream
 - Scalar inputs: AXI4-Lite interface (s_axilite)
 - Pointer to an Array:
 - AXI4 memory-mapped interface (m_axi) to access memory and
 - s_axilite interface to specify the base address to the memory address space.
 - Function Return: ap_return port added to the s_axilite interface
 - Arguments specified as hls::stream: default to AXI4-Stream interface (axis)
- HLS will produce an associated set of C driver files during the Export RTL process

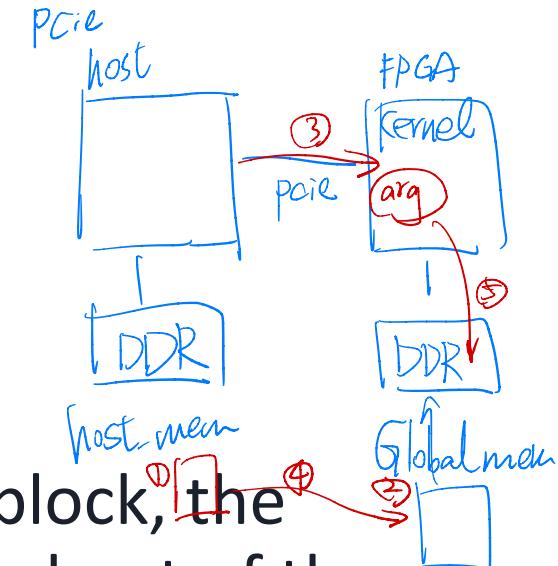


Block-Level Interface Protocol

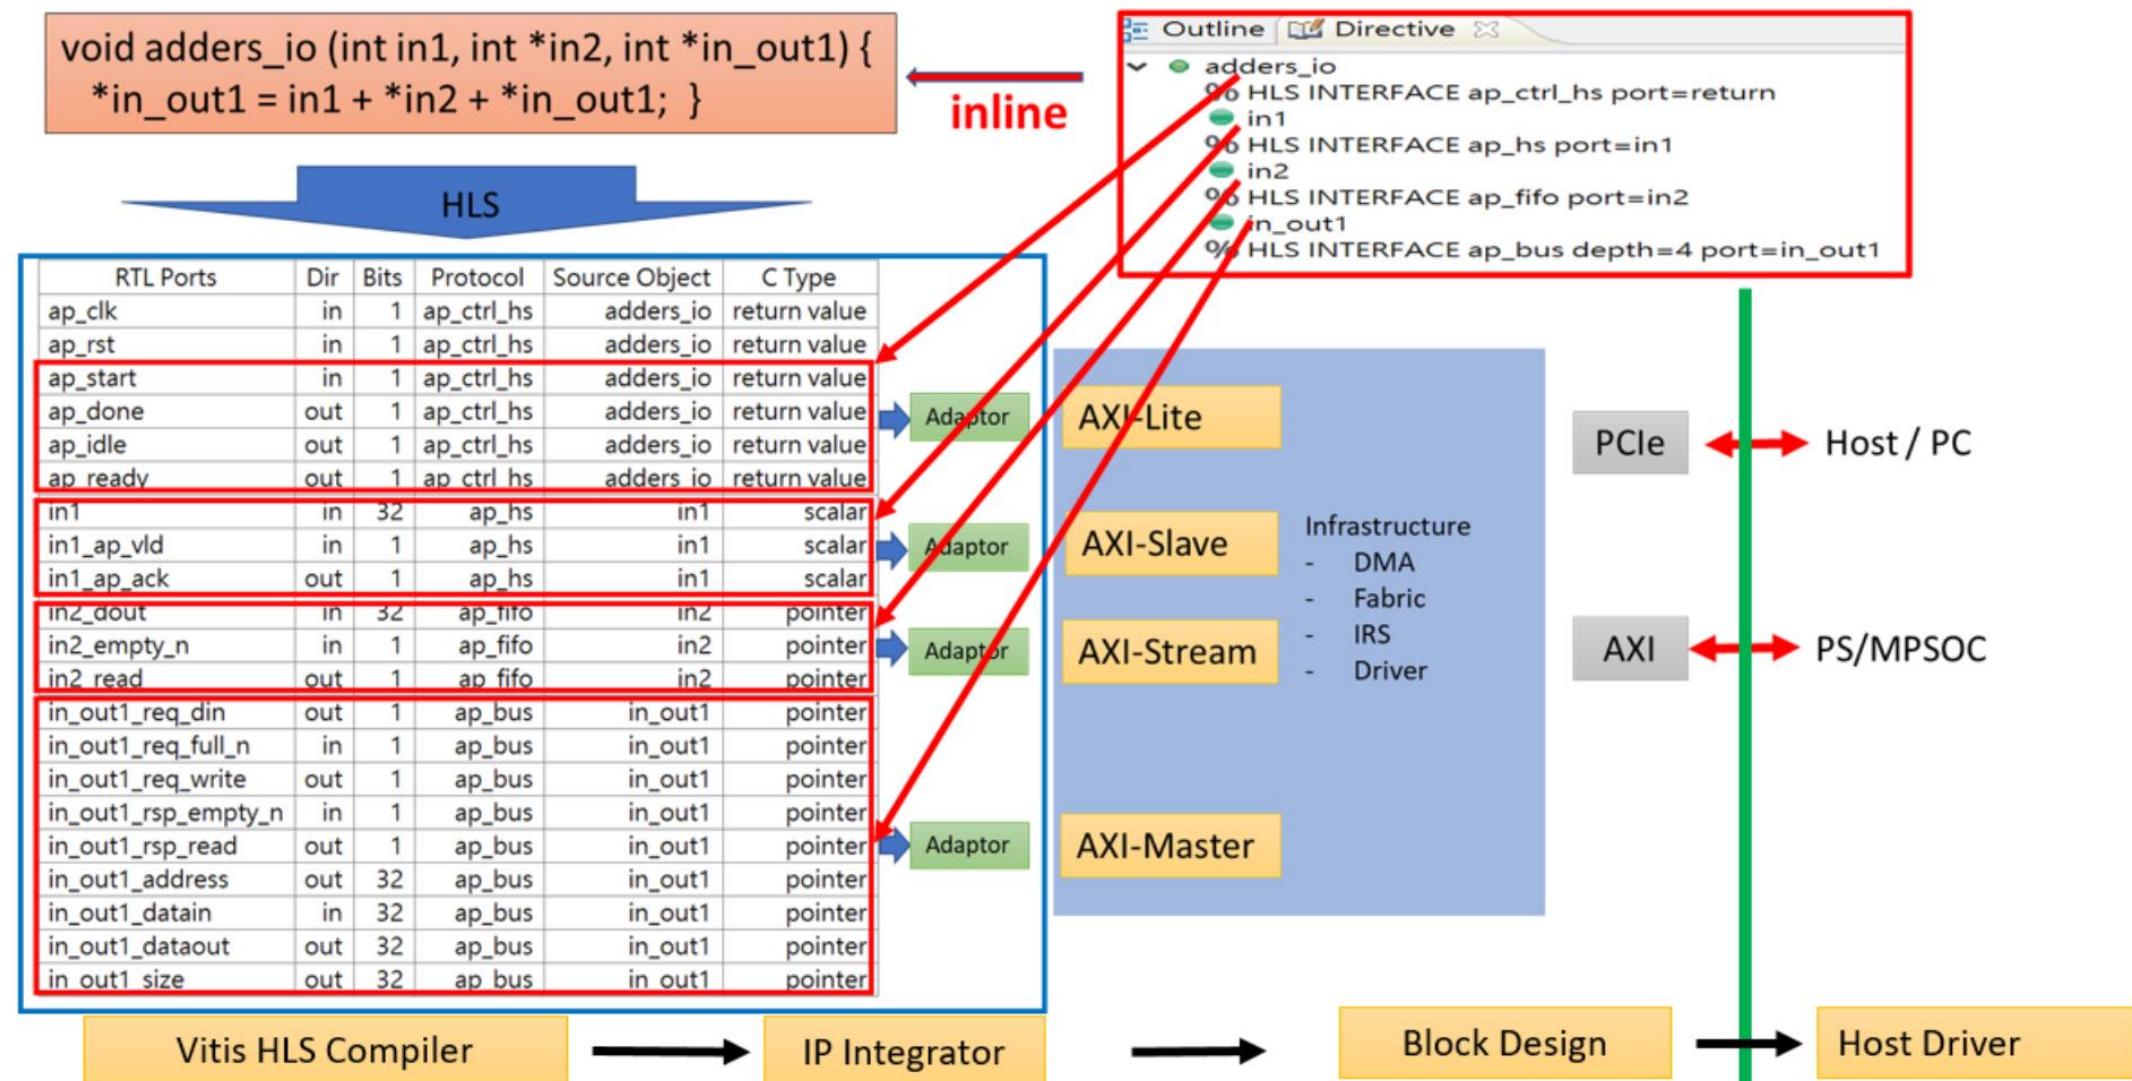


6	// 0x00 : Control signals
7	//
8	//
9	//
10	//
11	//
12	//
	bit 0 - ap_start (Read/Write/COH)
	bit 1 - ap_done (Read/COR)
	bit 2 - ap_idle (Read)
	bit 3 - ap_ready (Read)
	bit 7 - auto_restart (Read/Write)
	others - reserved

- Host application/driver controls kernel functions
- Block-level protocol (ap_ctrl_hs, ap_ctrl_chain) specifies:
 - When the design can start to perform the operation `ap_start`
 - When the operation ends `ap_done`
 - When the design is idle and ready for new inputs
- After the block-level protocol starts the operation of the block, the port-level I/O protocols are used to sequence data into and out of the block



Construction Flow for Host/Kernel Communication



Block Level Interface Protocol

↓ AP_CTRL_HS (Default: Sequential Mode)

AP_CTRL_CHAIN (Pipeline Mode)

AP_CTRL_NONE (Free-running Mode)

AP_CTRL_HS (Default)

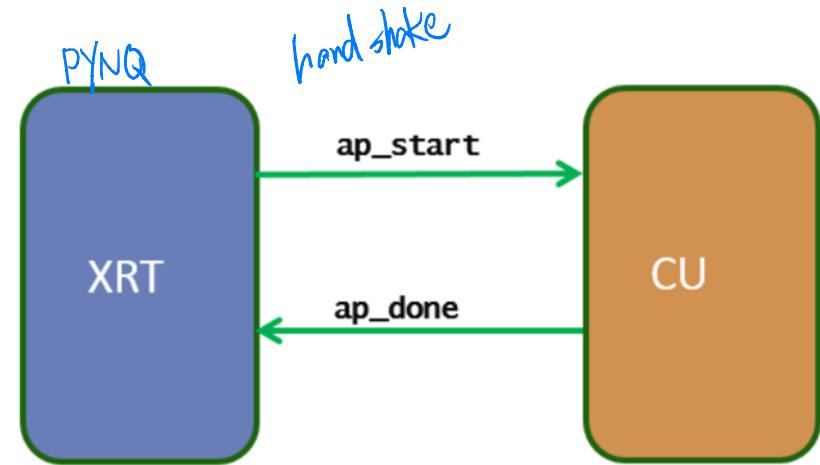
Block Level Protocol Specified on port: return

```
// kernel without chain  
void krnl_simple_mmult(int* a, int* b, int* c, int* d, int* output, int dim) {  
....  
#pragma HLS INTERFACE ap_ctrl_hs port = return  
no return port
```

(pipe → AP_CTRL_chain)

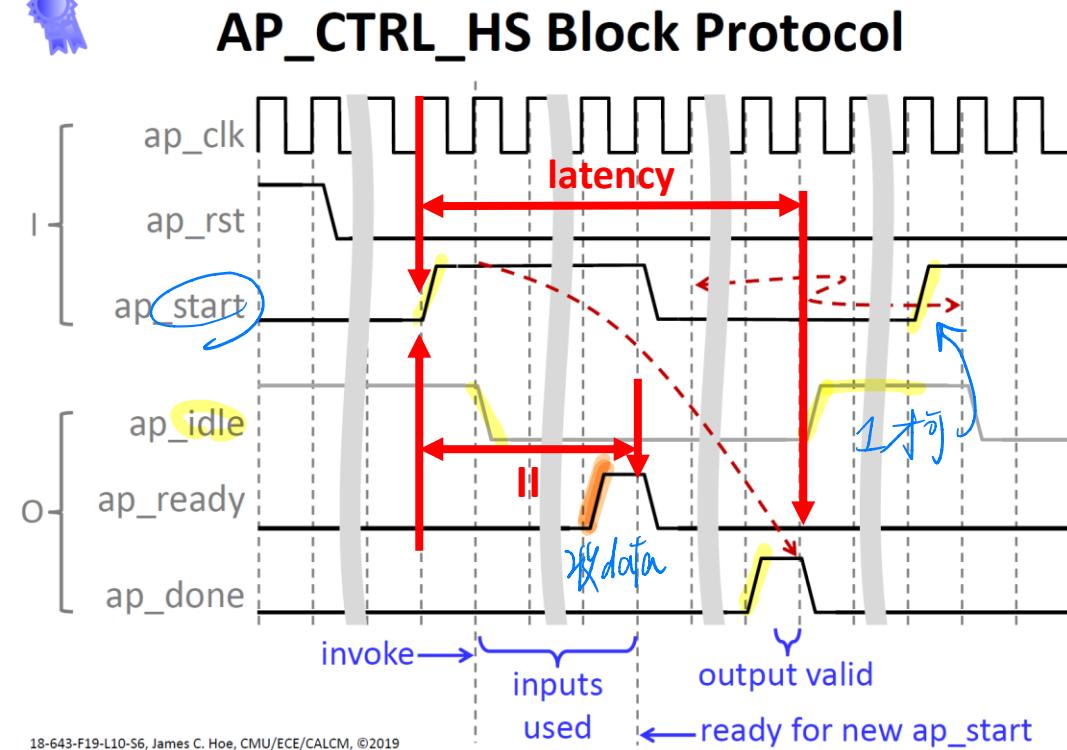
AP_CTRL_HS (Sequential Executed Kernel) non-pipe

- Host and Kernel Synchronization by
 - ap_start
 - ap_done
- Kernel can only be restarted (ap_start), after it completes the current execution (ap_done) *→ ap done 不可送 start (non-pipe)*
- Serving one execution request a time



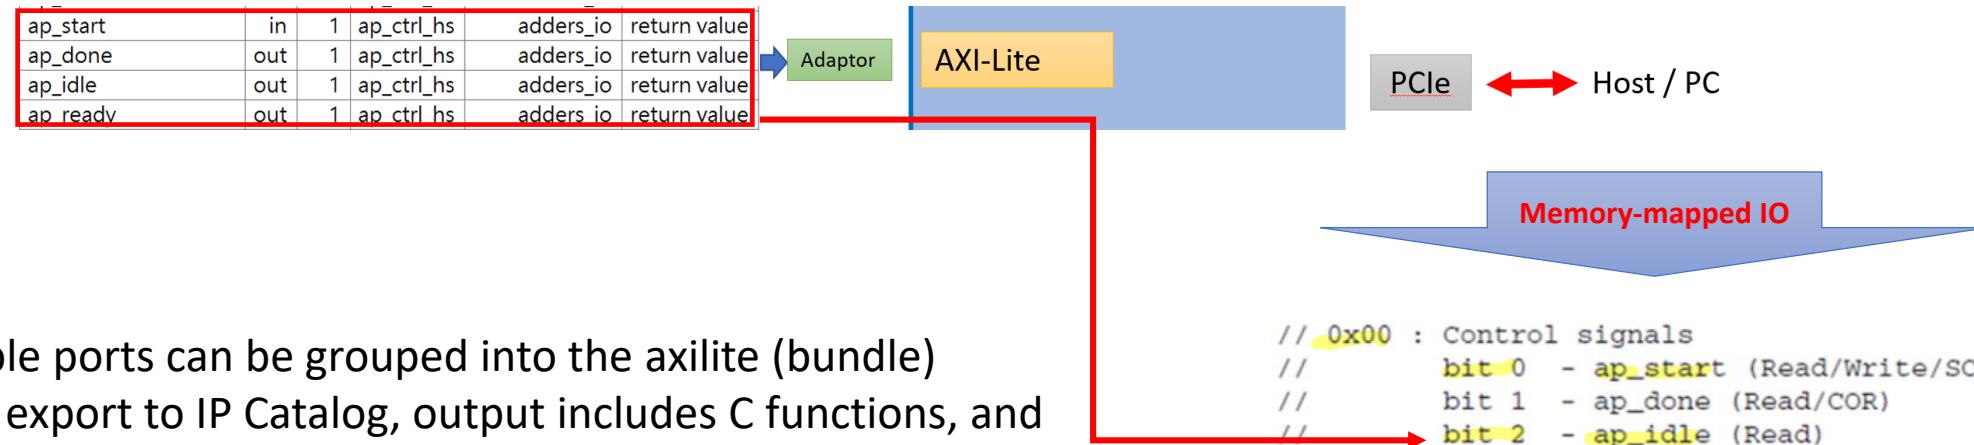
AP_CTRL_HS Protocol

- ap_start (i): set 1 to start until ap_ready asserted.
- ap_ready (o): design is ready to accept new input
- ap_done (o): design completes all operation.
Indicates data on ap_return is valid
- ap_idle (o): indicate design is idle if high.
- (ap_return): return data
- Pipeline/Nonpipeline depends on ap_ready timing



18-643-F19-L10-S6, James C. Hoe, CMU/ECE/CALCM, ©2019

Host (PS/Processor) to control HLS block – s_axilite



- Multiple ports can be grouped into the axilite (bundle)
- When export to IP Catalog, output includes C functions, and header file for the use when code running on a processor.
- **To start the block operation, ap_start register set to 1**
- **The block start reading inputs => data must be ready before start**
- When the block completes operation, the ap_done, ap_idle, and ap_ready registers will be set by the hardware output ports. And PS/Processor can read from the registers

```
// 0x00 : Control signals
// bit 0 - ap_start (Read/Write/SC)
// bit 1 - ap_done (Read/COR)
// bit 2 - ap_idle (Read)
// bit 3 - ap_ready (Read)
// bit 7 - auto_restart (Read/Write)
// others - reserved
// 0x04 : Global Interrupt Enable Register
// bit 0 - Global Interrupt Enable (Read/Write)
// others - reserved
// 0x08 : IP Interrupt Enable Register (Read/Write)
// bit 0 - Channel 0 (ap_done)
// bit 1 - Channel 1 (ap_ready)
// 0x0c : IP Interrupt Status Register (Read/TOW)
// bit 0 - Channel 0 (ap_done)
// others - reserved
```

0x10

AP_CTRL_HS : HOST/XRT

- The HOST/XRT driver writes a 1 in ap_start to start the kernel
- The HOST/XRT driver waits for ap_done asserted by kernel (guaranteeing the output data is fully produced by the kernel).
- Repeat 1-2 for the next kernel execution

AP_CTRL_CHAIN - Advanced

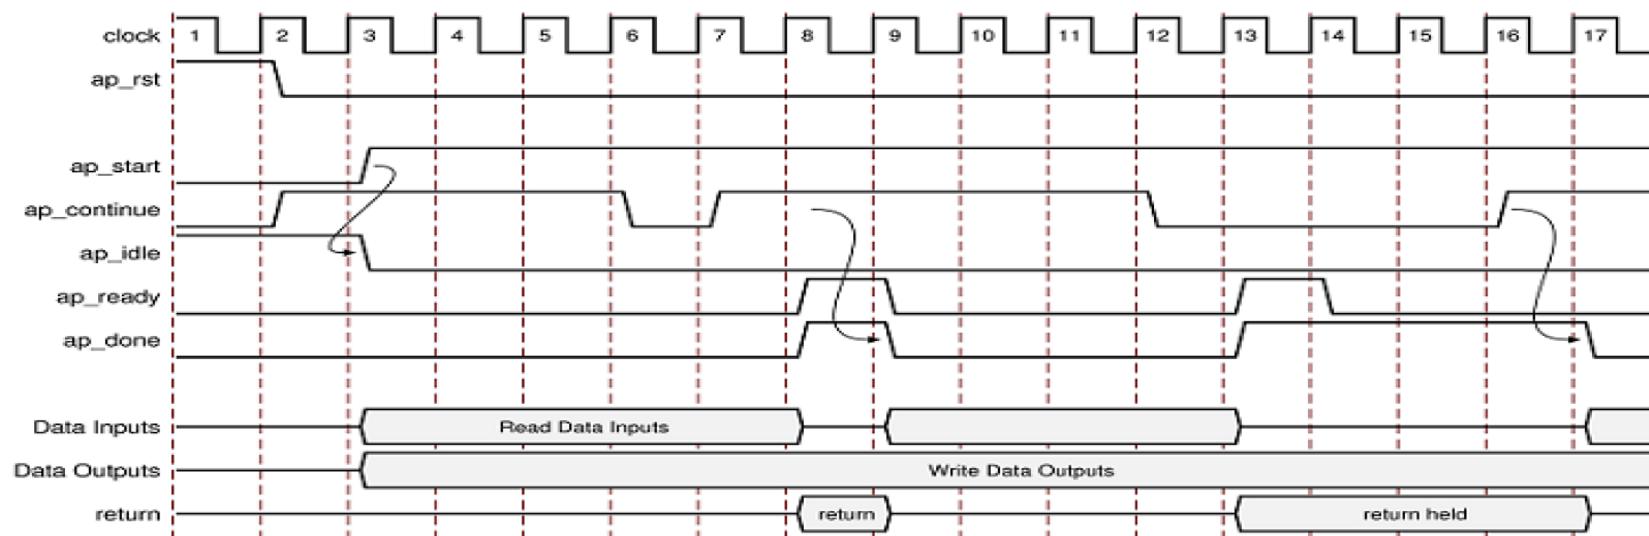
XAP_CTRL_CHAIN (Pipelined kernel)

For cascaded block, a **mechanism for consumer block to back-pressure on producer block**.

ap_continue – active high indicates the downstream block is ready to consumes for new data input.

- ap_continue is low, prevents the upstream block from generating additional data.
- down stream ap_ready can drive ap_continue port.
- ap_continues is High when ap_done is High, continue operating
- ap_continue Low, ap_done is High, stops operating, ap_done remains high, data remains valid on the ap_return port.

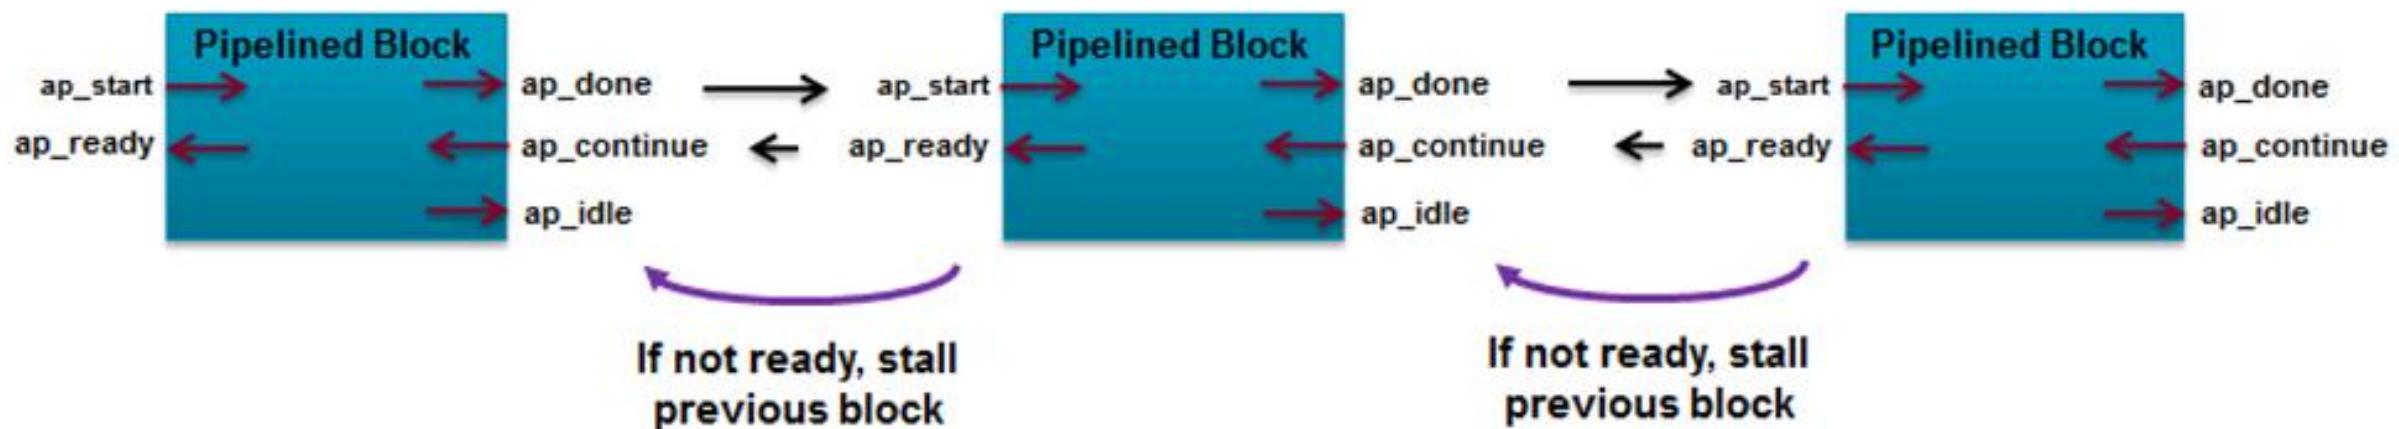
Figure 100: Behavior of ap_ctrl_chain Interface



AP_CTRL_CHAIN

> Protocol to support pipeline chains: ap_ctrl_chain

- >> Similar to ap_ctrl_hs but with additional signal ap_continue
- >> Allows blocks to be easily chained in a pipelined manner
- >> ap_ctrl_chain protocol provides back-pressure in systems

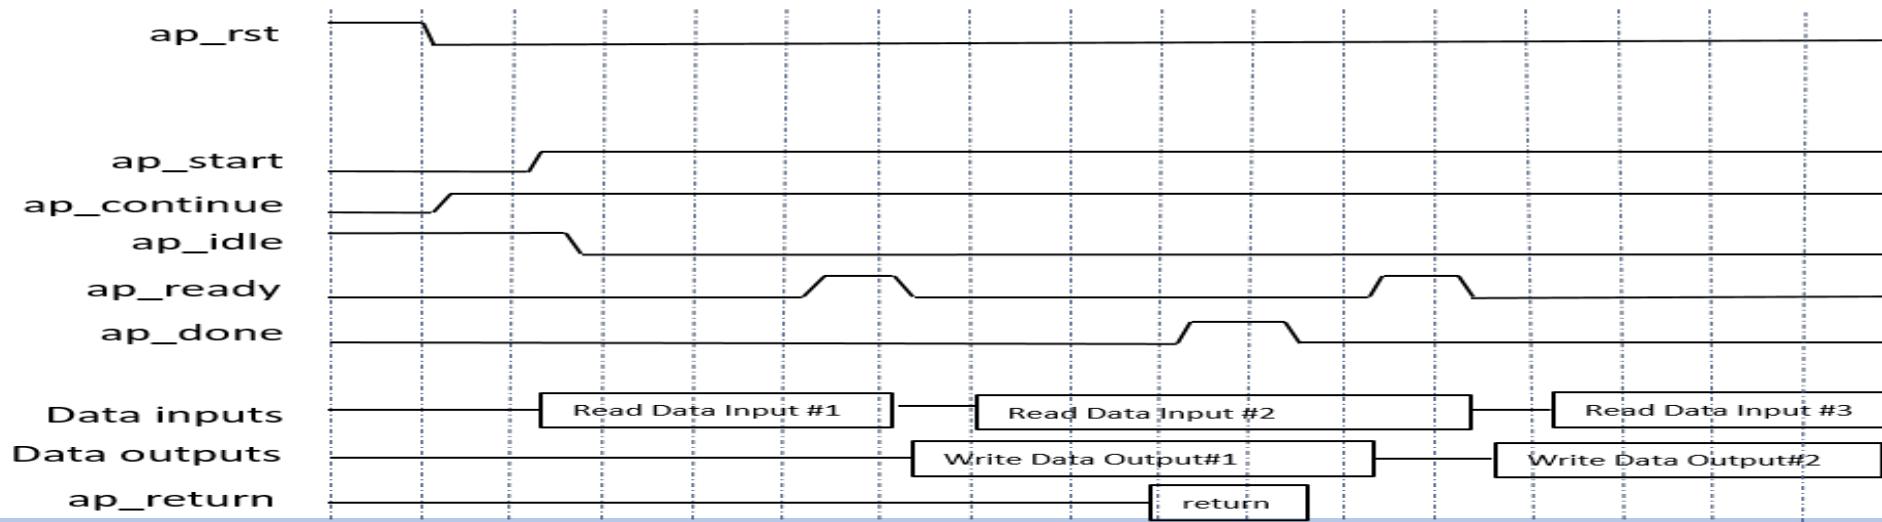
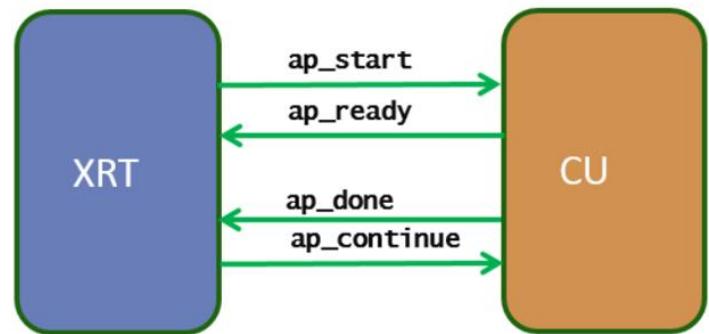


Host Control Pipeline Kernel Execution

1. Input Synchronization(ap_start, ap_ready)
 - XRT writes a 1 in ap_start to start the kernel
 - XRT waits for ap_ready
 - XRT write 1 in ap_start to start the kernel again
2. Output Synchronization (ap_done, ap_continue)
 - XRT waits for ap_done asserted by the kernel (output data is produced)
 - XRT write1 a 1 in ap_continue to keep kernel running.
3. The two processes (Input Synchronization, Output Synchronization) run asynchronously
4. Enqueue multiple requests, and data buffers ahead of time, e.g. clEnqueueMigrateMemObjects

Table 9: Control Register Signals

Bit	Name	Description
0	ap_start	Asserted when the kernel can start processing data. Cleared on handshake with ap_done being asserted.
1	ap_done	Asserted when the kernel has completed operation. Cleared on read.
2	ap_idle	Asserted when the kernel is idle.
3	ap_ready	Asserted by the kernel when it is ready to accept the new data
4	ap_continue	Asserted by the XRT to allow kernel keep running
7	auto_restart	Used to enable automatic kernel restart as described in Working with Auto-Restarting Kernels .
31:5	Reserved	Reserved



Performance Difference between pipeline, non-pipeline

```
// kernel without chain
void krnl_simple_mmmt(int* a, int* b, int* c, int* d, int* output, int dim) {
...
#pragma HLS INTERFACE ap_ctrl_hs port = return
```

```
// kernel with chain
void krnl_chain_mmmt(int* a, int* b, int* c, int* d, int* output, int dim) {
...
#pragma HLS INTERFACE ap_ctrl_chain port = return
```

```
#pragma HLS DATAFLOW
mm2s(a, strm_a, strm_ctrl_trans1, strm_ctrl_trans2);
mmult(strm_a, b, strm_ctrl_trans2, strm_b, strm_ctrl_trans3);
mmult(strm_b, c, strm_ctrl_trans3, strm_c, strm_ctrl_trans4);
mmult(strm_c, d, strm_ctrl_trans4, strm_d, strm_ctrl_trans5);
s2mm(strm_d, output, strm_ctrl_trans5);
}
```

```
// host code
for (int i = 0; i < NUM_TIMES; i++) {
    krnl_chain_mmmt.setArg(0, buffer_in1[i]);
    krnl_chain_mmmt.setArg(1, buffer_in2[i]);
    krnl_chain_mmmt.setArg(2, buffer_in3[i]);
    krnl_chain_mmmt.setArg(3, buffer_in4[i]);
    krnl_chain_mmmt.setArg(4, buffer_output[i]);
    krnl_chain_mmmt.setArg(5, MAT_DIM));
}

// Copy input data to device global memory
q.enqueueMigrateMemObjects({buffer_in1[i], buffer_in2[i], buffer_in3[i],
    buffer_in4[i]}, 0 /* 0 means from host*/);
q.enqueueTask(krnl_chain_mmmt);
}
```

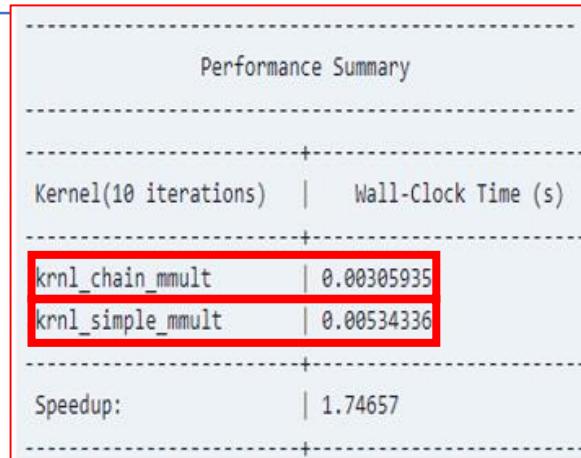
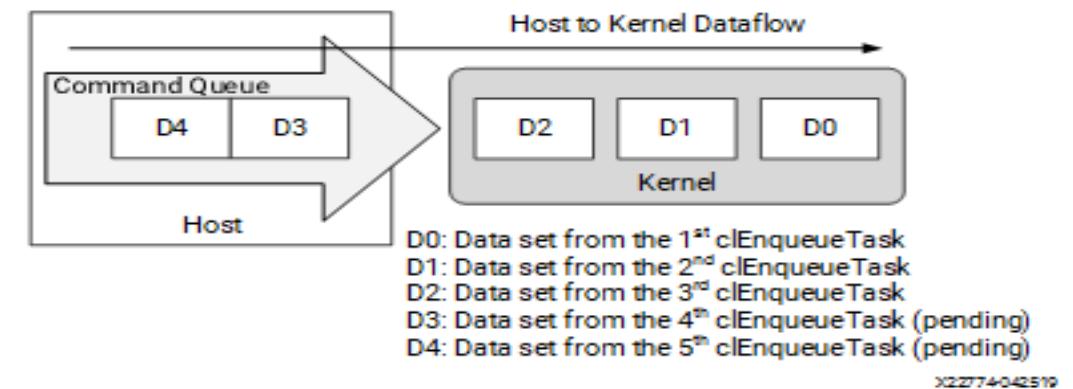


Figure: Host to Kernel Dataflow



AP_CTRL_NONE – For Dataflow

AP_CTRL_NONE (Continuously Running Kernel)

Host communicates with kernel through stream - Kernel starts execution when the data is available at its input

- Only when it has no memory-mapped input and output.
- No need to start the kernel by `clEnqueueTask` from the host.
- Useful for data-driven designs with streaming data coming from and going to the I/O pins (Ethernet, SerDes) of the FPGA, or streamed from or to a different kernel (kernel-to-kernel streaming).
- Don't use `clSetKernelArg` to pass a scalar argument to `ap_ctrl_none` kernel; only use `xclRegWrite` (API implemented in 2019.2) API

> Requirement: Manually verify the RTL

- » Without block level handshakes autosim cannot verify the design
 - Will only work in simple combo and II=1 cases
 - Handshakes are required to know when to sample output signals

It is recommended to leave the default and use Block Level Handshakes

Port Level – AXI Interface

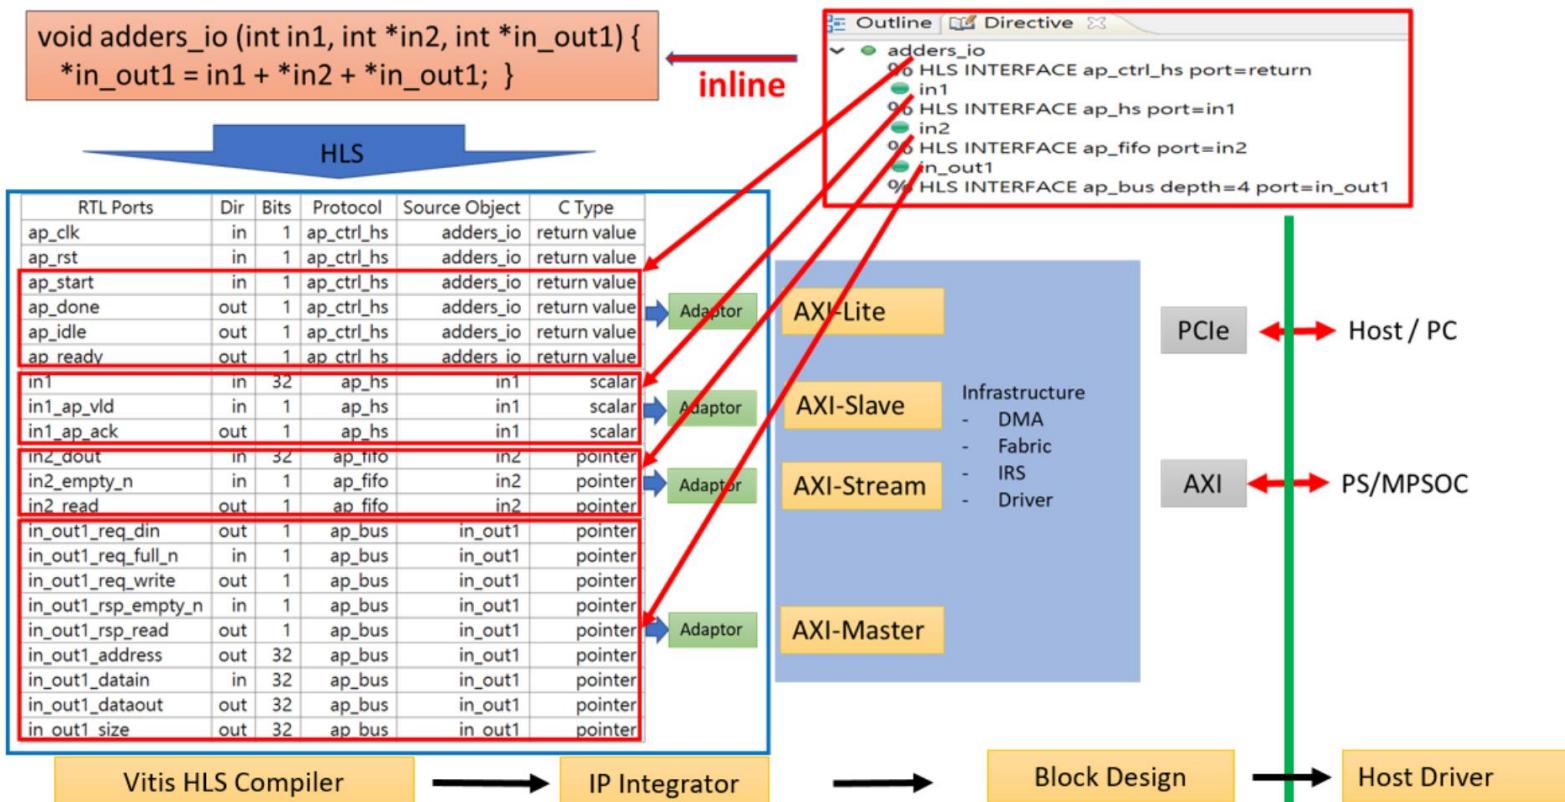
s_axilite,

m_axi

axis

axis, s_axilite, m_axi

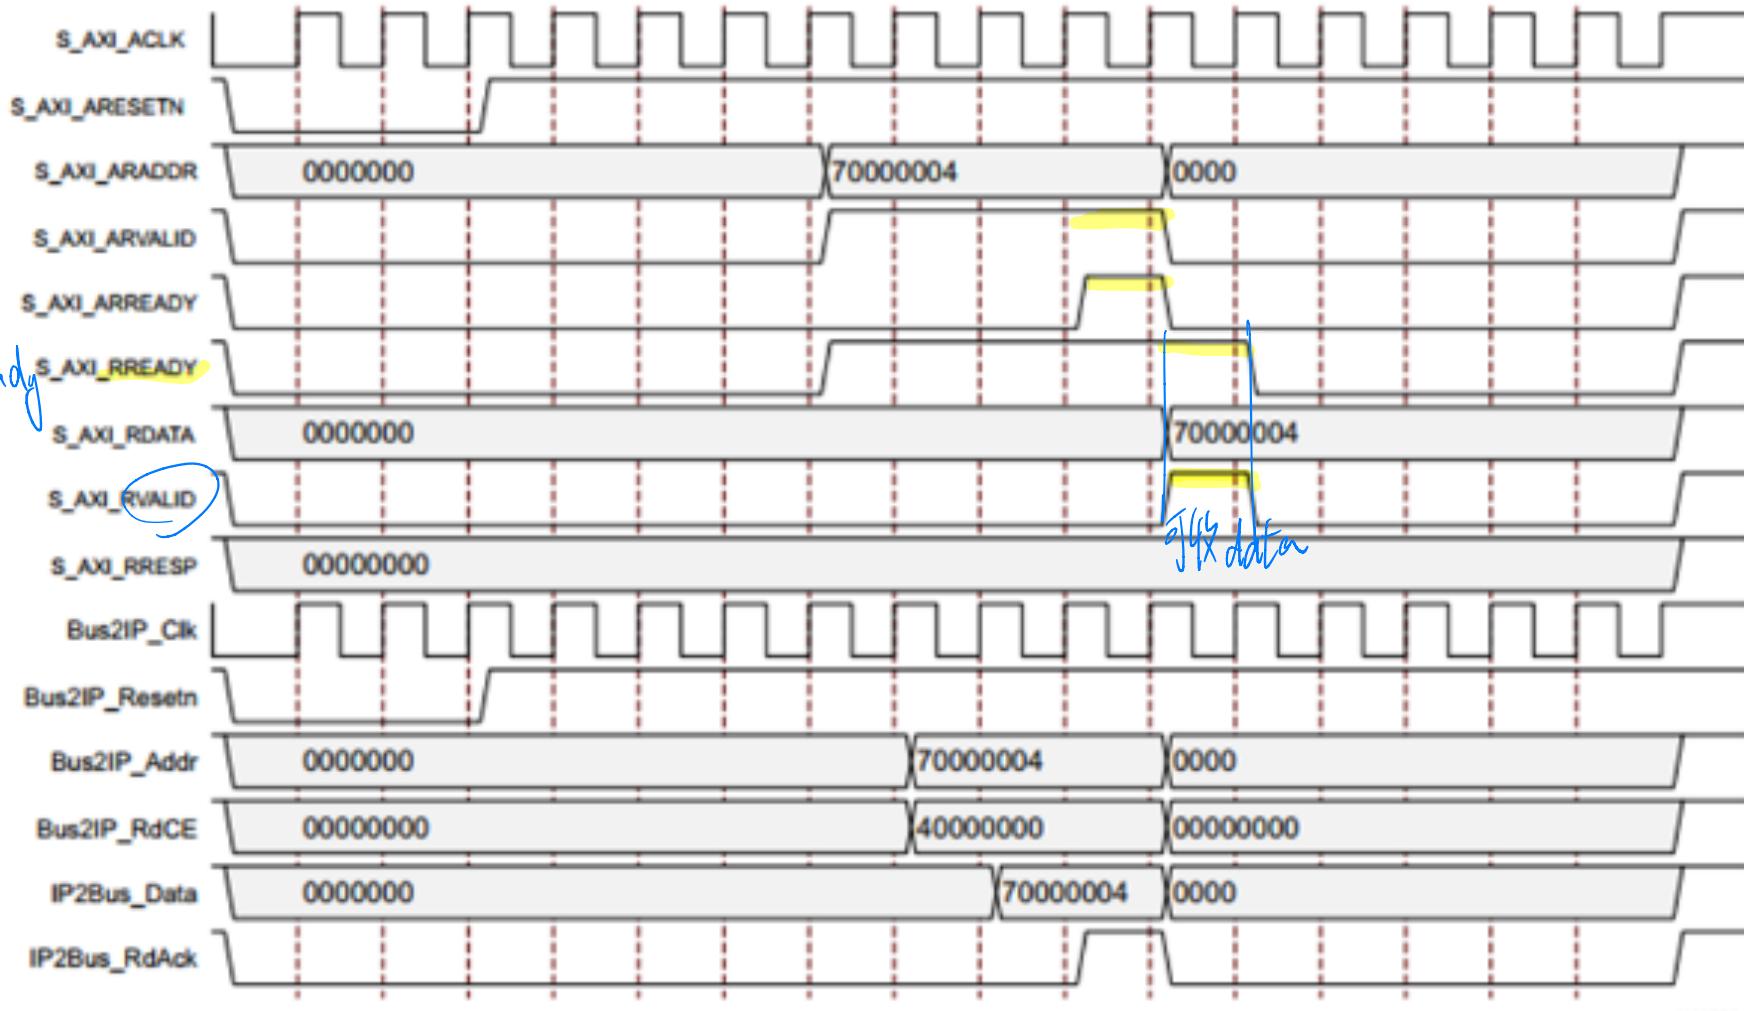
- s_axilite - AXI4-Lite – group multiple argument into the same AXI4-Lite
- m_axi - AXI4 Master – array or pointer, group multiple ports (bundle)
- axis - AXI4-Stream – for input or output but not io, array



Port Level – S-AXILITE

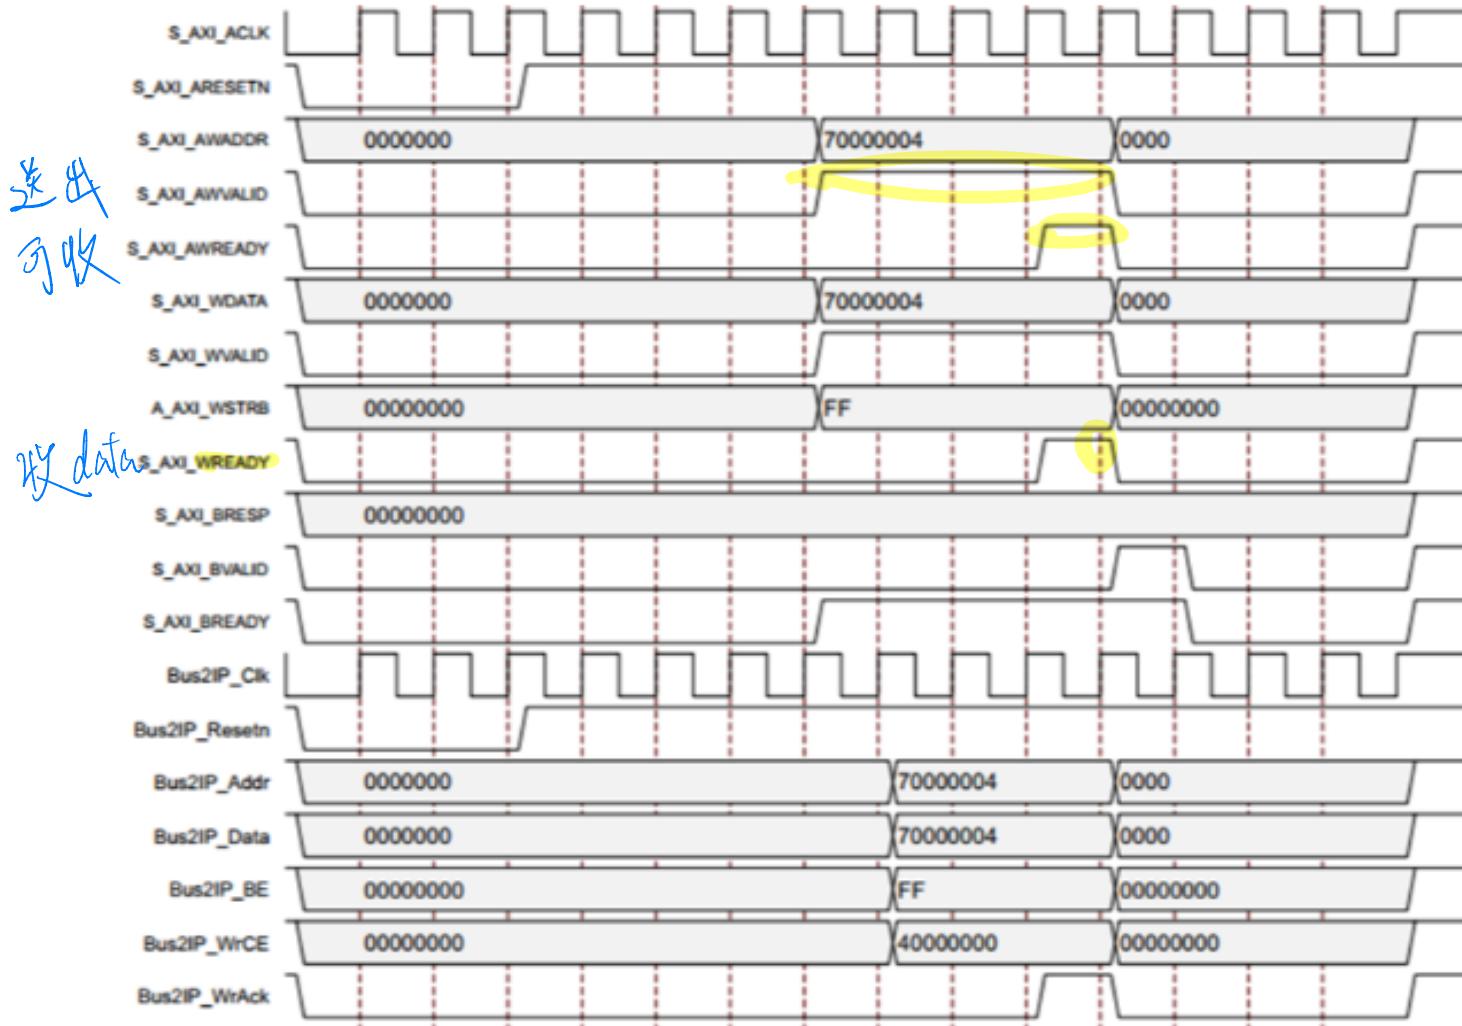
AXI4-Lite Read Transaction

AW (addr write)
AR (addr read)
R
Burst



AXI4-Lite Write Transaction

non-pipe

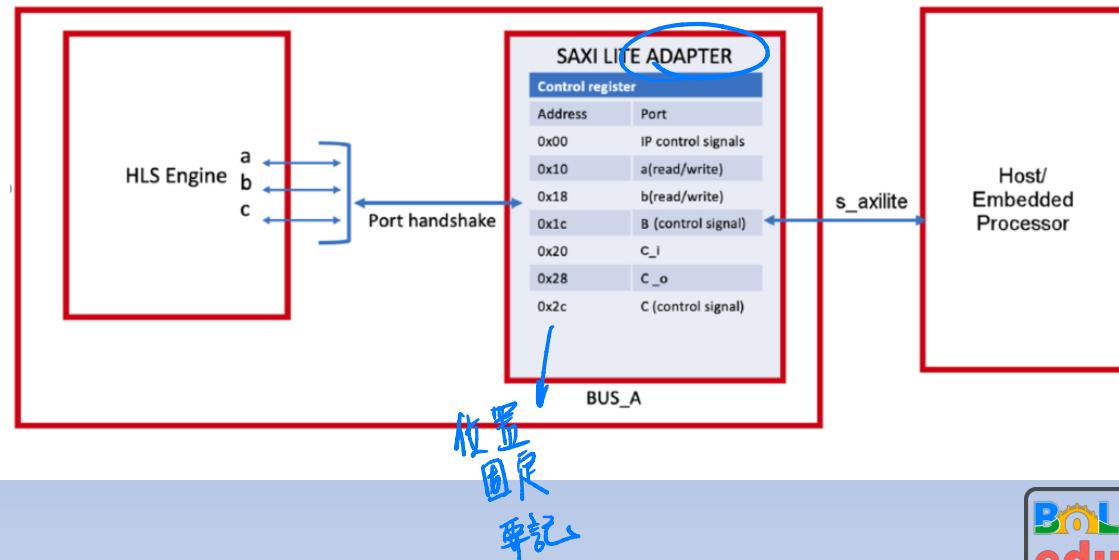


AXI-Lite Interface (s_axilite)

- Allow design controlled by a CPU
- **Port return:** specify the block-level protocol. (s_axilite: default)
- Bundle: Group multiple arguments, or separate axilite
- Offset: specified starting address, or tool automatically assigns
- Output C driver files for use with code running on a processor
- Address 0x0000-0x000F for block I/O protocol and interrupt control
- When the AXI4-lite interface is selected
 - Default mode for input ports is ap_none
 - Default mode for output port is ap_vld
 - Default mode for function return port is ap_ctrl_hs
- Controlling Hardware – _hw.h
 - Start block operation ap_start register set to 1
 - Block complete ap_done, ap_idle, and ap_ready set by hardware

```
#include <stdio.h>
void example(char *a, char *b, char *c)
{
    #pragma HLS INTERFACE s_axilite  port=return bundle=BUS_A
    #pragma HLS INTERFACE ap_ctrl_hs port=return bundle=BUS_A
    #pragma HLS INTERFACE s_axilite  port=a   bundle=BUS_A
    #pragma HLS INTERFACE s_axilite  port=b   bundle=BUS_A
    #pragma HLS INTERFACE s_axilite  port=c   bundle=BUS_A
    #pragma HLS INTERFACE ap_vld   port=b

    *c += *a + *b;
}
```



S_AXILITE and Block Level Protocol

- Port return: specify the block-level protocol. The interface pragma is specified as

#pragma HLS INTERFACE s_axilite port=return bundle=BUS_A

It is the default block-level protocol. 固定

- You can also assign the block control protocol to the interface. As an example

#pragma HLS INTERFACE s_axilite port=return bundle=BUS_A

#pragma HLS INTERFACE ap_ctrl_hs port=return bundle=BUS_A AXI Lite register

- In the Control Register Map, HLS reserve address 0x00 – 0x0C for the block level protocols and interrupt controls.

- The Control register (0x00) contains ap_start (bit-0), ap_done (bit-1), ap_ready (bit-3) and ap_idle (bit_2); in the case of ap_ctrl_chain, it contains ap_continue. These signals are accessed through the s_axilite adapter.

S-AXILITE and Port-Level Protocol

```
#include <stdio.h>
void example(char *a, char *b, char *c)
{
    #pragma HLS INTERFACE s_axilite port=return bundle=BUS_A
    #pragma HLS INTERFACE ap_ctrl_hs port=return bundle=BUS_A
    #pragma HLS INTERFACE s_axilite port=a      bundle=BUS_A
    #pragma HLS INTERFACE s_axilite port=b      bundle=BUS_A
    #pragma HLS INTERFACE s_axilite port=c      bundle=BUS_A
    #pragma HLS INTERFACE ap_vld      port=b

    *c += *a + *b;
}
```

```
#pragma HLS INTERFACE s_axilite port=b bundle=BUS_A
#pragma HLS INTERFACE ap_vld port=b

// 0x18 : Data signal of b
//           bit 7~0 - b[7:0] (Read/Write)
//           others - reserved
// 0x1c : Control signal of b
//           bit 0 - b_ap_vld (Read/Write/SC)
//           others - reserved
```

```
#pragma HLS INTERFACE s_axilite port=a bundle=BUS_A
```

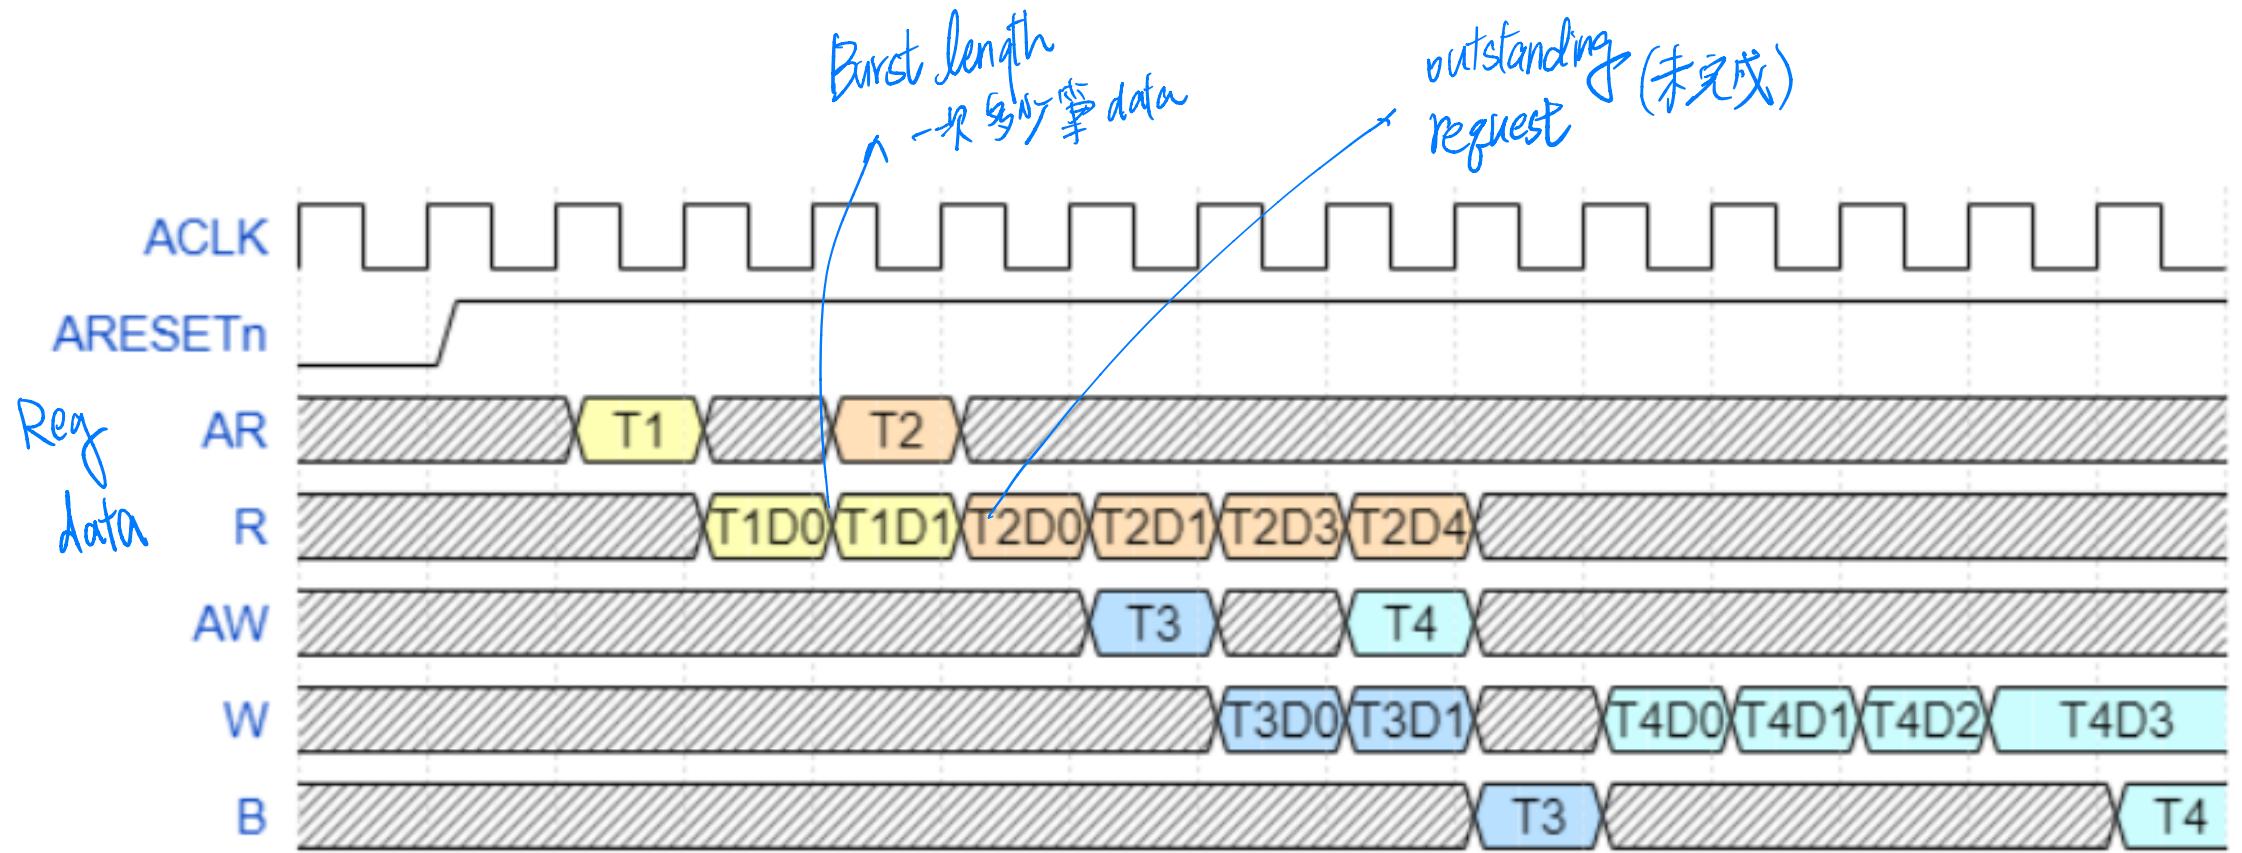
```
// 0x10 : Data signal of a
//
//           bit 7~0 - a[7:0] (Read/Write)
//           others - reserved
```

```
#pragma HLS INTERFACE s_axilite port=c bundle=BUS_A
```

```
// 0x20 : Data signal of c_i
//
//           bit 7~0 - c_i[7:0] (Read/Write)
//           others - reserved
// 0x24 : reserved
// 0x28 : Data signal of c_o
//
//           bit 7~0 - c_o[7:0] (Read)
//           others - reserved
// 0x2c : Control signal of c_o
//
//           bit 0 - c_o_ap_vld (Read/COR)
//           others - reserved
```

Port Level – AXI4_Master (M_AXI)

AXI Master Transaction Waveform



Example of multiple transactions: T1 (arlen=1), T2 (arlen=3) read transactions and T3 (awlen=1), T4 (awlen=3) write transactions.

Single v.s. Burst Transfer

- Pointer and array argument (default to m_axi)
- Offset = slave, transfer address is defined by axilite.
- Default alignment is set to 64 bytes
- Maximum read/write burst length is set to 16 by default

Single Transfer

Burst Transfer

- memcpy
- For-loop + PIPELINE
 - Pipeline the loop
 - Access address **in increasing order**
 - Do not place accesses inside a **conditional statement**
 - For nested loops, **do not flatten loops**, because this inhibits the burst operation
 - **Only one read and one write is allowed in a for loop** unless the ports are bundled in different AXI ports.

Burst
i=5

```
void bus (int *d) {  
    static int acc = 0;  
    acc += *d;  
    *d = acc;  
}
```

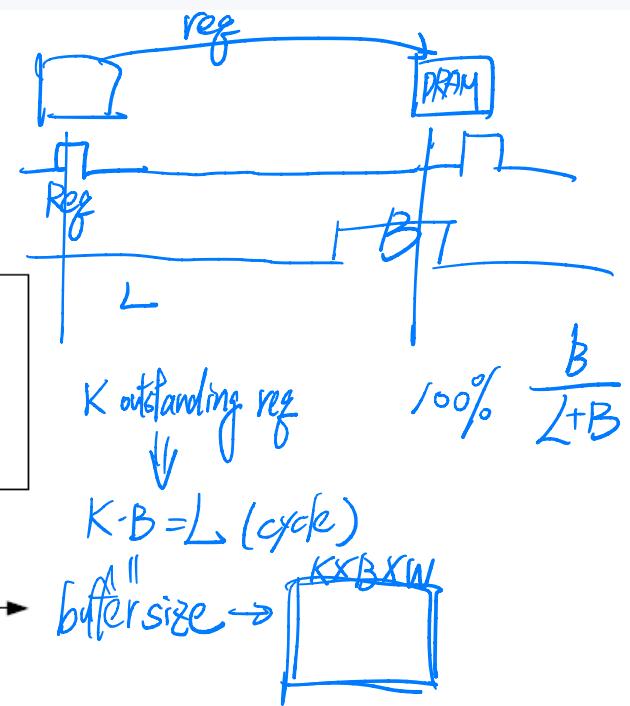
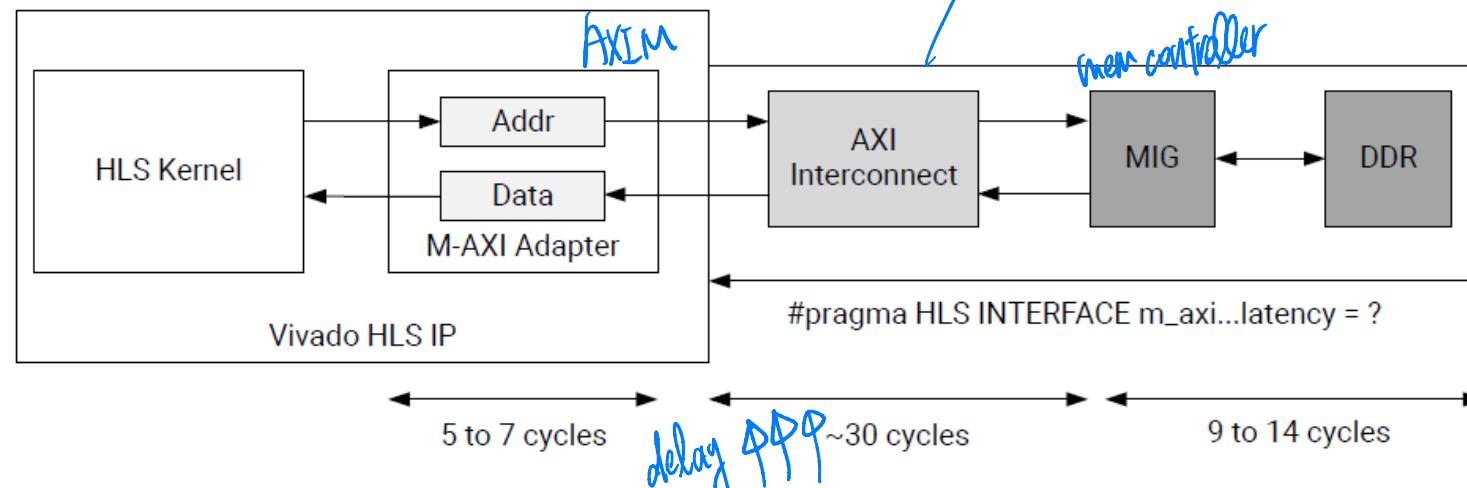
```
void example(volatile int *a){  
    #pragma HLS INTERFACE m_axi      depth=50   port=a  
    #pragma HLS INTERFACE s_axilite port=return  
    int i, buff[5];  
  
    //memcpy creates a burst access to memory  
    ① memcpy(buff, (const int*) a, 50*sizeof(int));  
  
    for(i=0; i < 50; i++){  
        #pragma HLS PIPELINE  
        buff[i] = buff[i] + 100; }  
  
    // for loop + PIPELINE  
    ② for(i=0; i < 50; i++){  
        #pragma HLS PIPELINE  
        a[i] = buff[i]; }  
    }  
    ↓寫出 data
```

AXI Master Characteristics

bus width burst length
 $\leq 4KB$

- Max burst-length ≤ 256 (limited by AXI bus protocol ARLEN[7:0])
 8-bit
- Max transfer size $\leq 4KB$
- Do not cross 4KB boundary
- Bus width – power of 2, between 32 bits and 512 bits
 - What if not matches actual data size, e.g. video data RGB 24-bit?

Latency & Bus utilization



- **depth**: Specifies the maximum number of samples for the test bench to process.
- **latency**: Specifies the expected latency of AXI4 interface.
- **max_read_burst_length, max_write_burst_length**: Specifies the maximum number of data values read during a burst transfer.
- **num_read_outstanding, num_write_outstanding**: Specifies how many read requests can be made to the AXI4 bus without a response before the design stalls. It needs storage to hold num_read_outstanding requests
Buffer size = num_read_outstanding * max_read_burst_length * word_size.

Question: With 128-bit bus width, what is the max burst length to specify?

Best Practices for Designing with M_AXI Interfaces

https://docs.xilinx.com/r/en-US/ug1399-vitis-hls/Best-Practices-for-Designing-with-M_AXI-Interfaces

Experiment on AXI Burst Performance

- AXI burst performance affected by
 - Latency
 - Burst Length
 - Num_outstanding
 - Bus width
- Measure the time to read/write a buffer from DDR

```
// Data Width - 256
void test_kernel_maxi_256bit_1(int64_t buf_size, int direction, int64_t* perf,
ap_int<256>* mem) {
// Data Width - 512
// void test_kernel_maxi_512bit_1(int64_t buf_size, int direction, int64_t* perf,
ap_int<512>* mem) {

#pragma HLS INTERFACE m_axi port = mem bundle = aximm0 \
    num_write_outstanding = 4 \
    max_write_burst_length = 4 \
    num_read_outstanding = 4 \
    max_read_burst_length = 4 \
    offset = slave
#pragma HLS DATAFLOW
hls::stream<int64_t> cmd;
testKernelProc(mem, buf_size, direction, cmd, 32);
perfCounterProc(cmd, perf, direction, 4, 4);
}
```

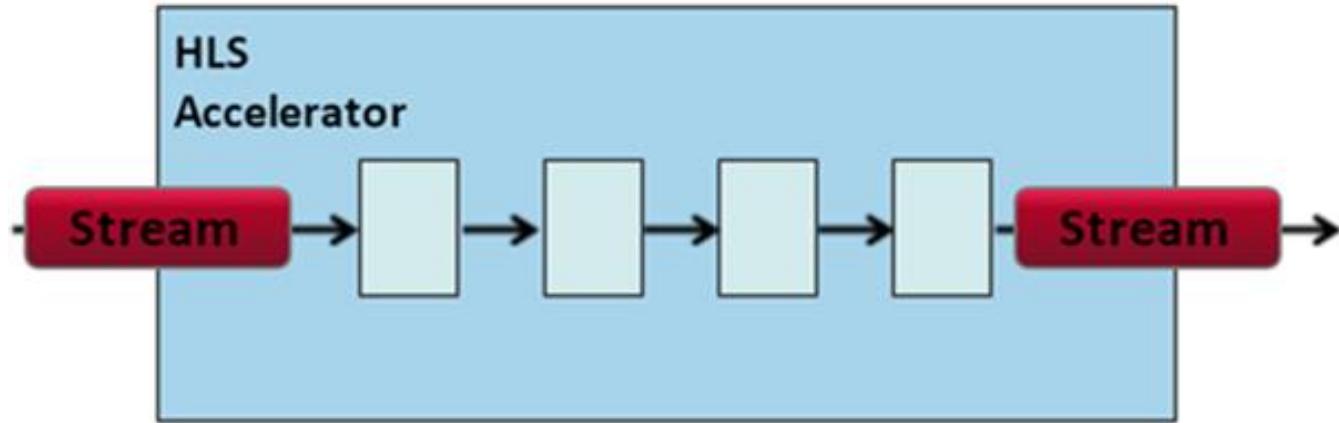
Kernel->AXI Burst READ performance

Data Width = 512 burst_length = 4 num_outstanding = 4 buffer_size = 16.00 MB throughput = 3.92988 GB/sec
Data Width = 512 burst_length = 16 num_outstanding = 4 buffer_size = 16.00 MB throughput = 13.1114 GB/sec
Data Width = 512 burst_length = 32 num_outstanding = 4 buffer_size = 16.00 MB throughput = 16.8218 GB/sec
Data Width = 512 burst_length = 4 num_outstanding = 32 buffer_size = 16.00 MB throughput = 16.8222 GB/sec
Data Width = 512 burst_length = 16 num_outstanding = 32 buffer_size = 16.00 MB throughput = 16.8295 GB/sec
Data Width = 512 burst_length = 32 num_outstanding = 32 buffer_size = 16.00 MB throughput = 16.8219 GB/sec

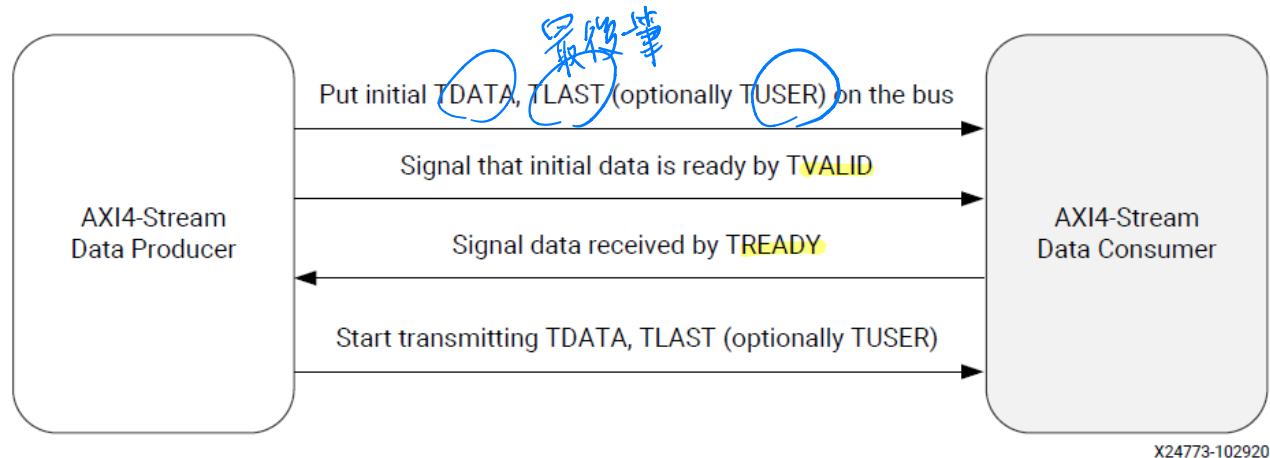
Port Level – AXI4_STREAM (AXIS)

AXI4-Stream - AXIS

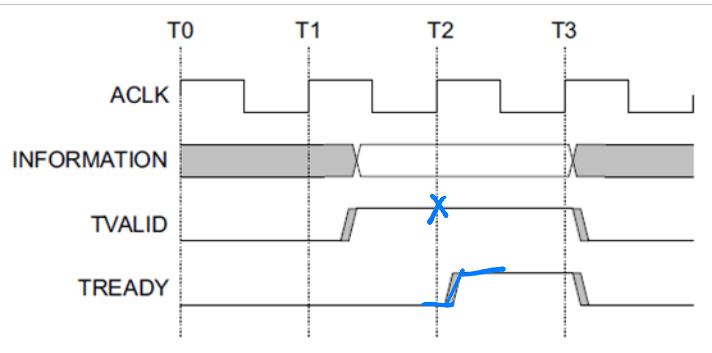
- Specify on input or output only , not on input/output – array, pointer
- Use mostly in dataflow between functions
- Multiple variables group into same stream (struct, DATA_PACK)



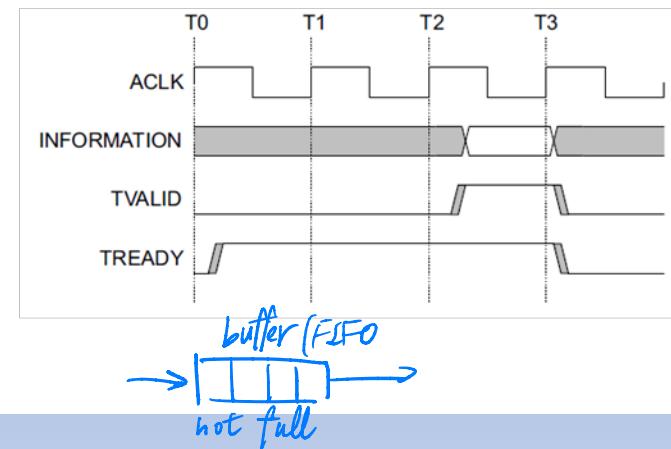
AXI4-Stream Transfer Protocol



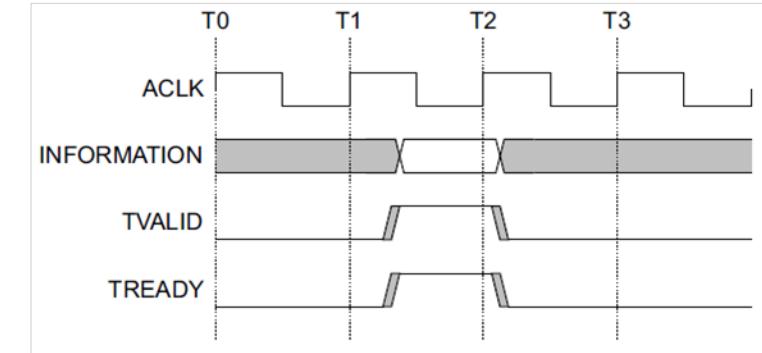
TVALID asserted before TREADY



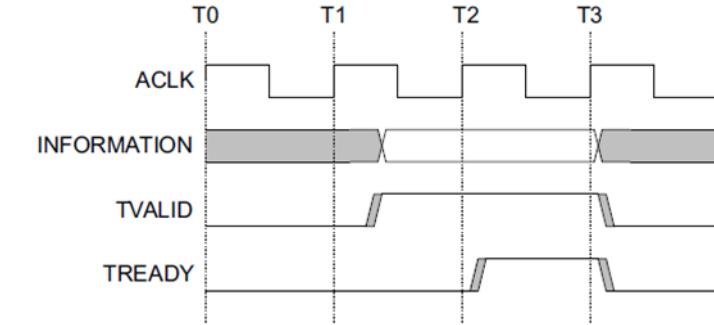
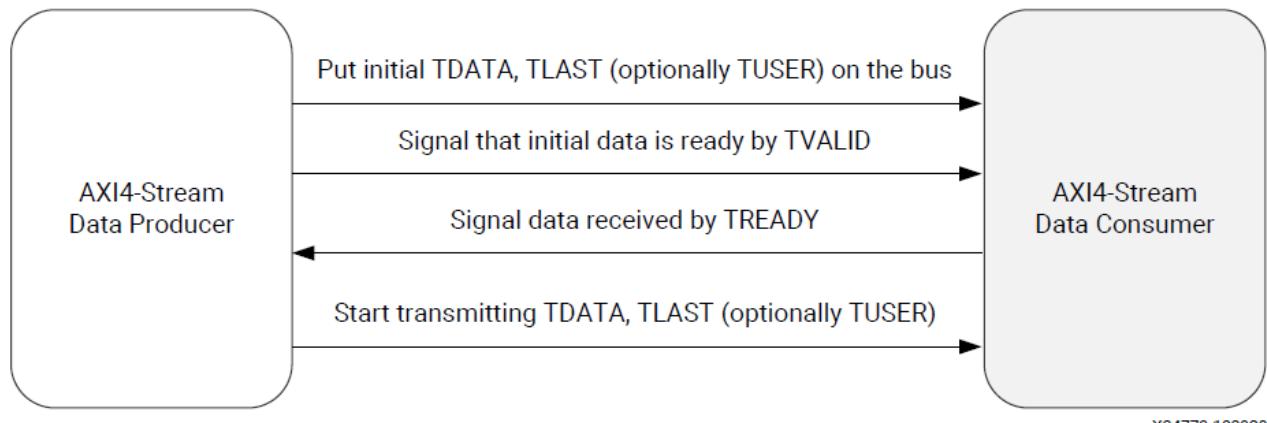
TREADY asserted before TVALID



TVALID and TREADY asserted simultaneously



AXI4-Stream – HLS::AXIS



```
template <typename T, size_t WUser, size_t WId, size_t WDest> struct axis {  
    static constexpr size_t NewWUser = (WUser == 0) ? 1 : WUser;  
    static constexpr size_t NewWId = (WId == 0) ? 1 : WId;  
    static constexpr size_t NewWDest = (WDest == 0) ? 1 : WDest;  
  
    T data;  
    ap_uint<bytewidth<T>> keep;  
    ap_uint<bytewidth<T>> strb;  
    ap_uint<NewWUser> user;  
    ap_uint<1> last;  
    ap_uint<NewWId> id;  
    ap_uint<NewWDest> dest;  
};
```

for Stream Router

```
ap_axis<Wdata, WUser, WId, WDest>  
hls::axis<ap_int<WData>, WUser, WId, WDest>
```

```
ap_axiu<WData, WUser, WId, WDest>  
hls::axis<ap_uint<WData>, WUser, WId, WDest>
```

Stream for Dataflow : `hls::stream<>`



- Include `<hls_stream.h>`
- **`hls::stream<Type, Depth>`**
 - Type:
 - C++ native data type
 - HLS arbitrary precision type, e.g. `ap_int<>`
 - User-defined struct containing above types
 - Depth: depth of FIFO for co-simulation verification
- Used for top-level function arguments, and between functions
- Top interface can be
 - FIFO interface: `ap_fifo` (default) - support non-blocking behavior
 - Handshake interface: `ap_hs`
 - AXI4-Stream : `axis`
- Inside function – FIFO with depth = 2 (`#pragma HLS STREAM depth = <int>`), note: depth **specify actual resource allocation**.
- Use passed-by-reference to pass streams into and out of functions
- Only in C++ based designs

Stream Example

> Create using hls::stream or define the hls namespace

```
#include <ap_int.h>
#include <hls_stream.h>

typedef ap_uint<128> uint128_t;           // 128-bit user defined type

hls::stream<uint128_t> my_wide_stream; // A stream declaration
```

```
#include <ap_int.h>
#include <hls_stream.h>
using namespace hls;                      // Use hls namespace

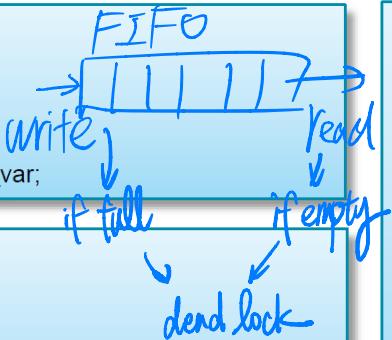
typedef ap_uint<128> uint128_t;           // 128-bit user defined type

stream<uint128_t> my_wide_stream; // hls:: no longer required
```

> Blocking and Non-Block accesses supported

```
// Blocking Write
hls::stream<int> my_stream;

int src_var = 42;
my_stream.write(src_var);
// OR use: my_stream << src_var;
```



```
// Blocking Read
hls::stream<int> my_stream;

int dst_var;
my_stream.read(dst_var);
// OR use: dst_var = my_stream.read();
```

```
hls::stream<int> my_stream;

int src_var = 42;
bool stream_full;

// Non-Blocking Write
if (my_stream.write_nb(src_var)) {
    // Perform standard operations
} else {
    // Write did not happen
}

// Full test
stream_full = my_stream.full();
```

```
hls::stream<int> my_stream;

int dst_var;
bool stream_empty;

// Non-Blocking Read
if (my_stream.read_nb(dst_var)) {
    // Perform standard operations
} else {
    // Read did not happen
}

// Empty test
fifo_empty = my_stream.empty();
```

Stream arrays and structs are not supported for RTL simulation at this time: must be verified manually.

e.g. hls::stream<uint8_t> chan[4]